

BUGS and R with NIMBLE (and automatic C++ compilation)

Perry de Valpine*

Daniel Turek

Chris Paciorek

Cliff Anderson-Bergman

Ras Bodik

Duncan Temple Lang

University of California

Berkeley

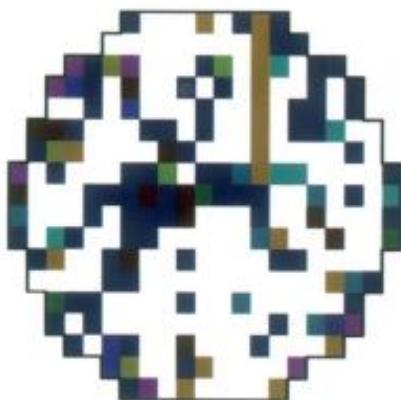
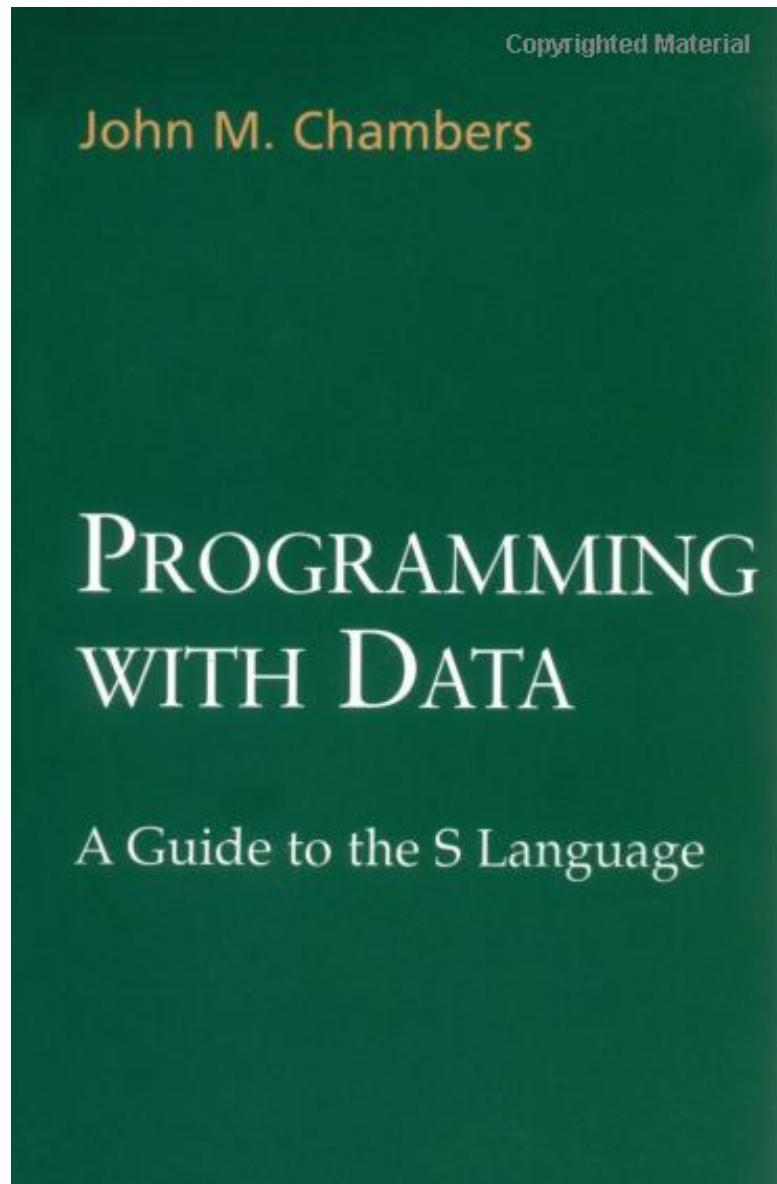
Davis



Funded by National Science Foundation
(Advances in Biological Informatics)

*Sabbatical host: Center for Math Research (CIMAT), Guanajuato, México

The core concept of R



(A lot of) Statistics = Models + Data

Packages reinvent the model part:

- Linear models: `lm(Y ~ A + B*X, ...)`
- Generalized linear mixed models: `glmer(Y~A+(1+x|B), ...)`
- Generalized additive models: `gam(Y ~ A + s(X), ...)`
- Time-series models
- Spatial models
- Survival models
- On and on...
- ...

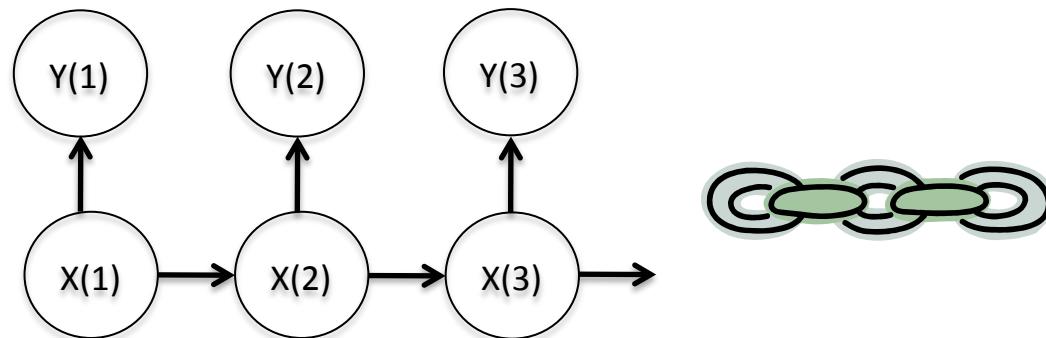
How about a language for models?

Languages for models

- BUGS
 - Bayesian inference Using Gibbs Sampling
 - WinBUGS, OpenBUGS, JAGS
 - motivated by MCMC
 - Most widely adopted
- Others: ADMB, PyMC, Stan, BLOG, Church
- Generally motivated by hierarchical models
 - Random effects and/or Bayesian analysis

Model languages have been tied to algorithms:

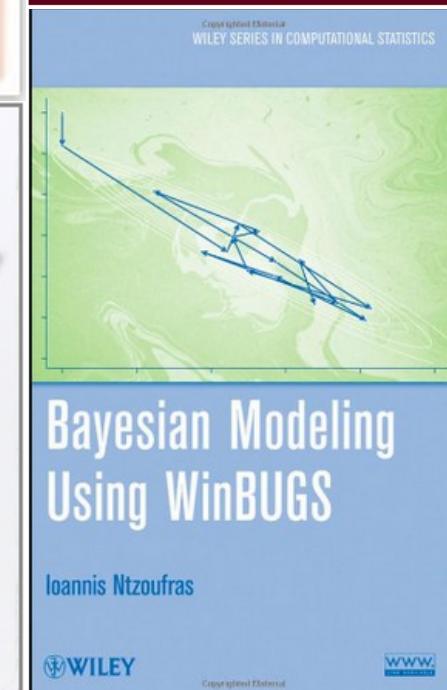
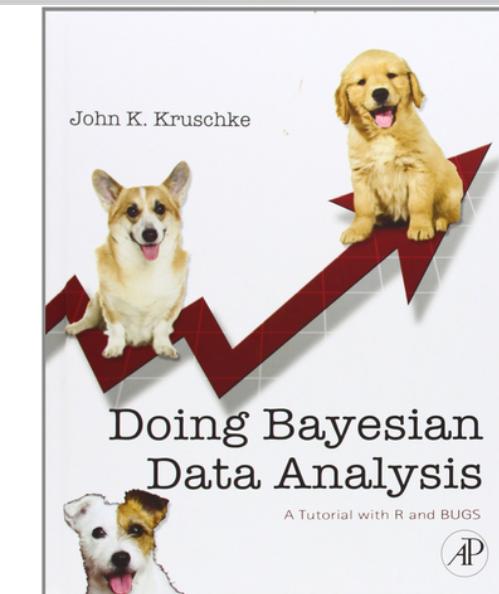
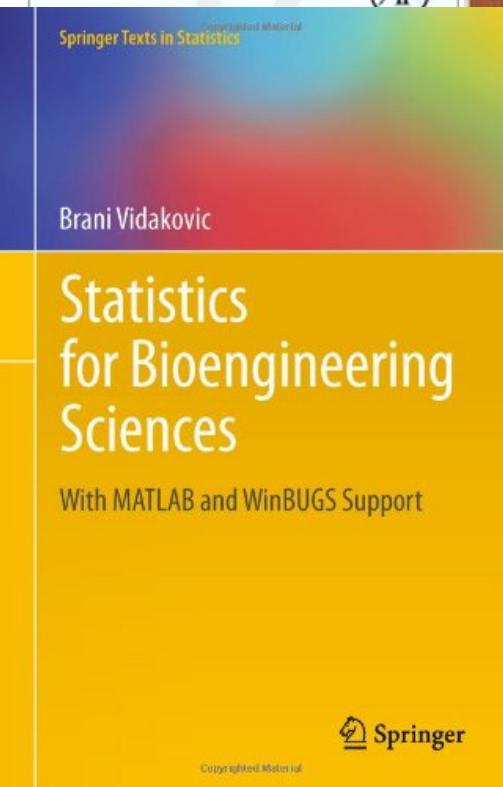
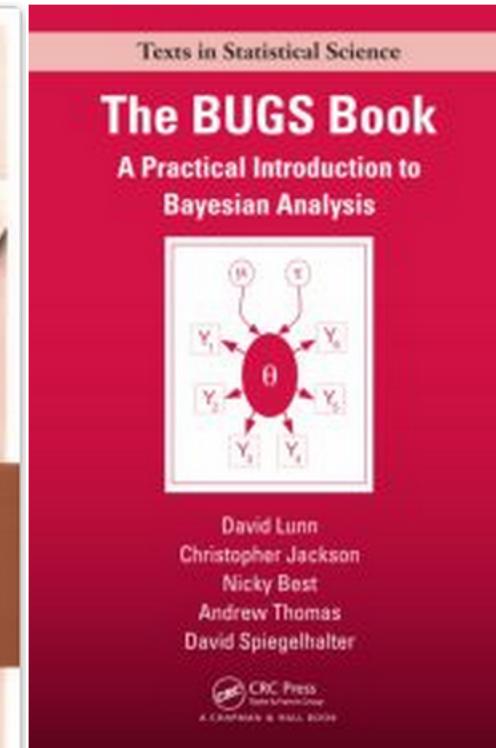
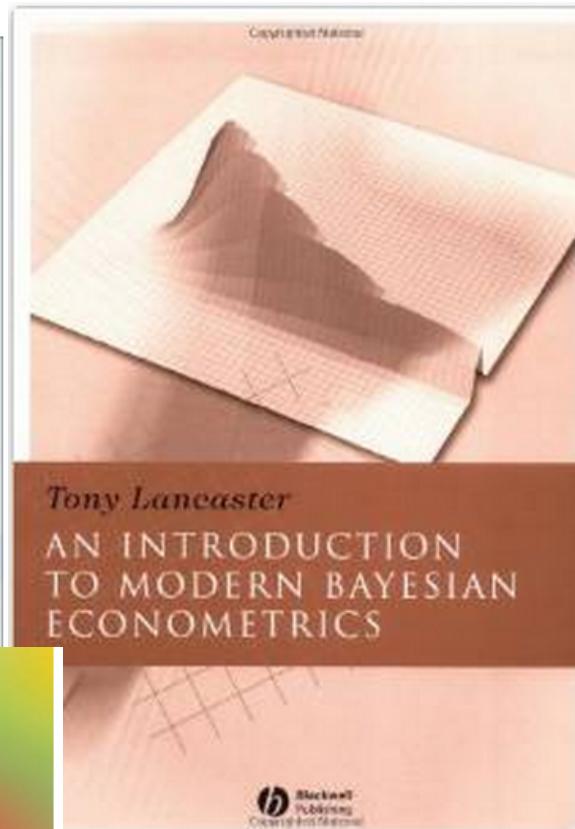
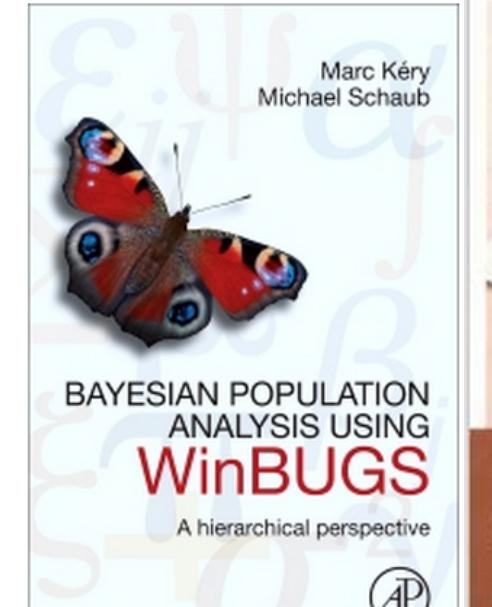
Model



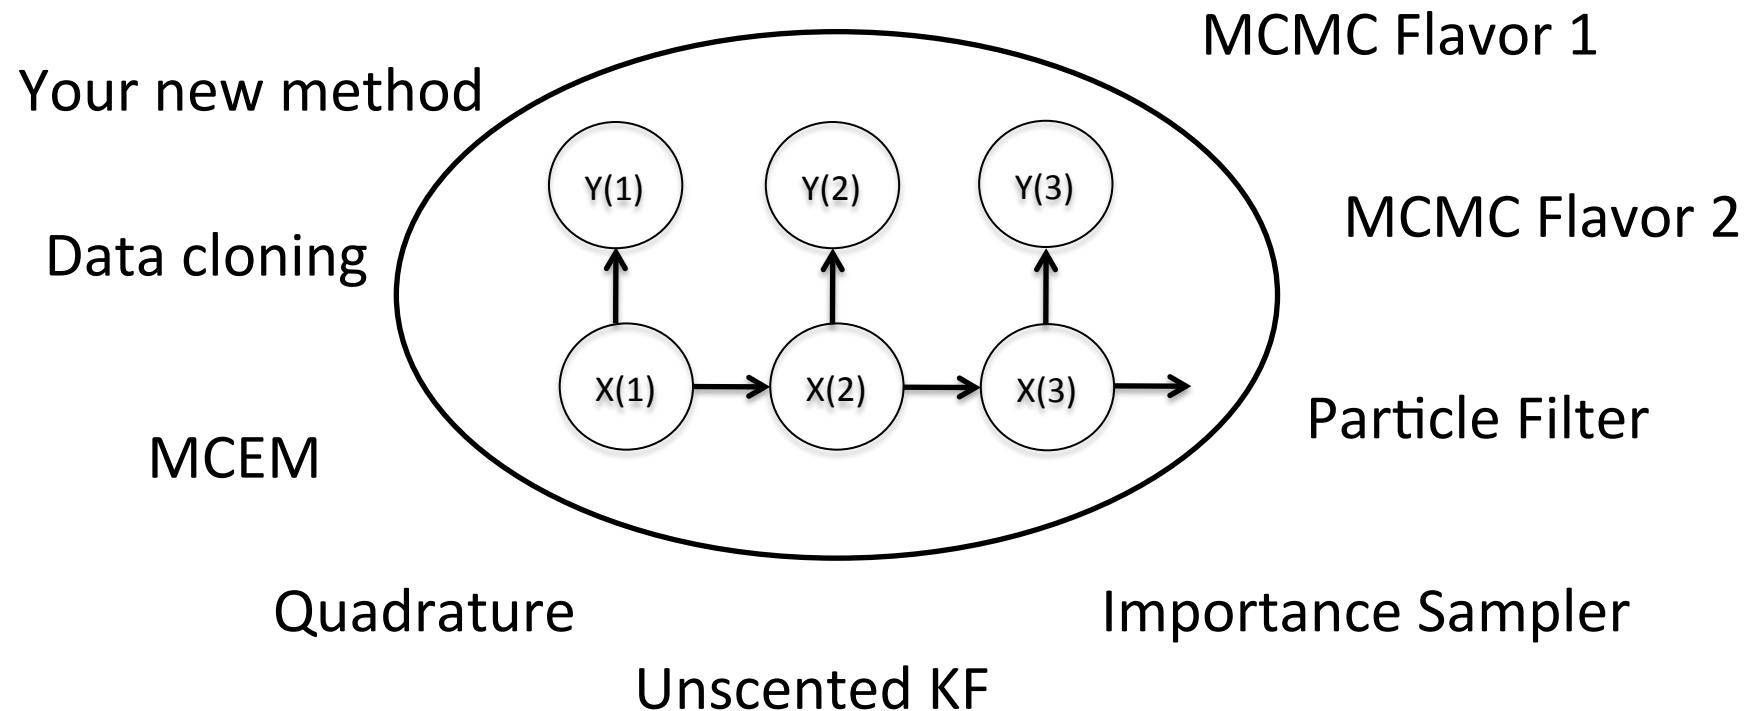
Algorithm



e.g. BUGS (WinBUGS, OpenBUGS, JAGS), ADMB, TMB, Stan

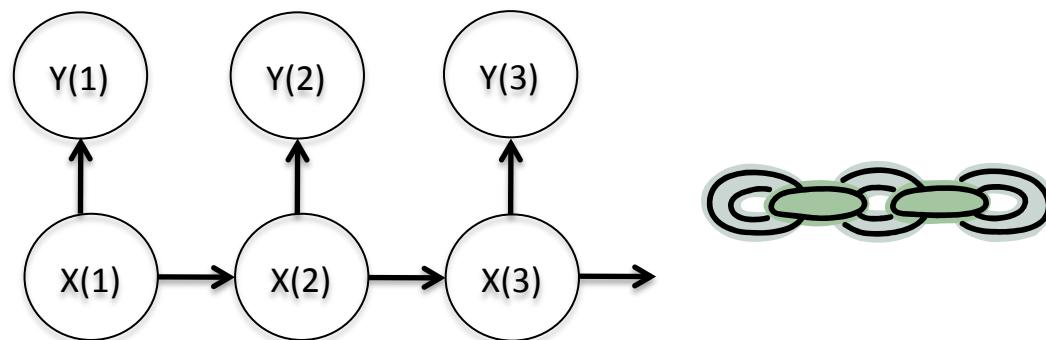


What do we want to do with a hierarchical model?



Model languages have been tied to algorithms:

Model



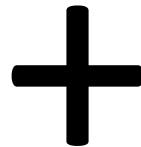
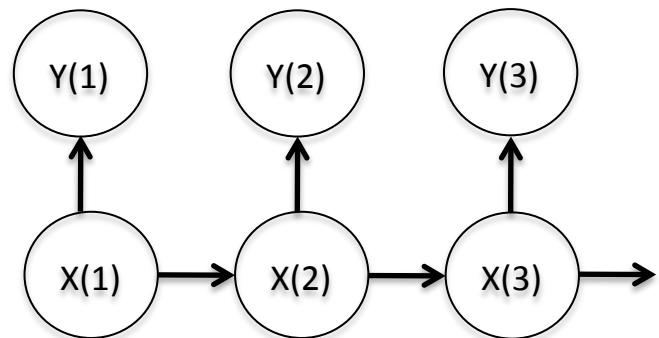
Algorithm



e.g. BUGS (WinBUGS, OpenBUGS, JAGS), ADMB, TMB, Stan

Goal of NIMBLE

Model (BUGS language)

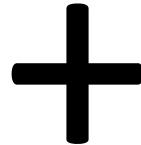
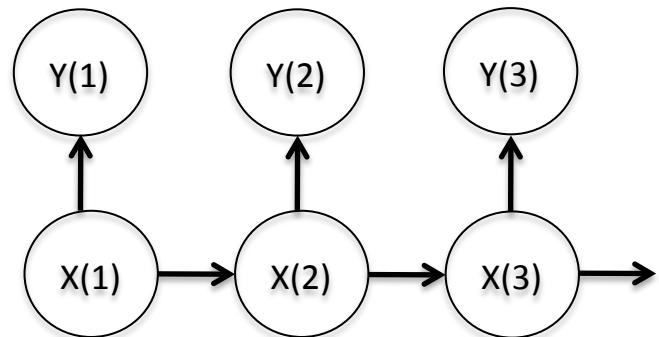


Algorithm Language



Goal of NIMBLE

Model (BUGS language)



Algorithm Language



MUST BE FAST = COMPILED

Numerical
Inference for statistical
Models using
Bayesian and
Likelihood
Estimation

VERSION 0.3-1

NIMBLE

1. BUGS language → R model object
2. NIMBLE programming language
 - ✧ Domain-specific language (DSL) embedded in R
 - ✧ Compilation via C++
3. Algorithm library: MLE, MCMC etc.

NIMBLE

1. BUGS language → R model object
2. NIMBLE programming language
 - ✧ Domain-specific language (DSL) embedded in R
 - ✧ Compilation via C++
3. Algorithm library: MLE, MCMC etc.

What NIMBLE does with BUGS models?

```
model {  
  for(i in 1:N) {  
    theta[i] ~ dgamma(alpha, beta)  
    lambda[i] <- theta[i] * tt[i]  
    x[i] ~ dpois(lambda[i])  
  }  
  alpha ~ dexp(1.0)  
  beta ~ dgamma(2, 2)  
}
```

(pump model example from BUGS)

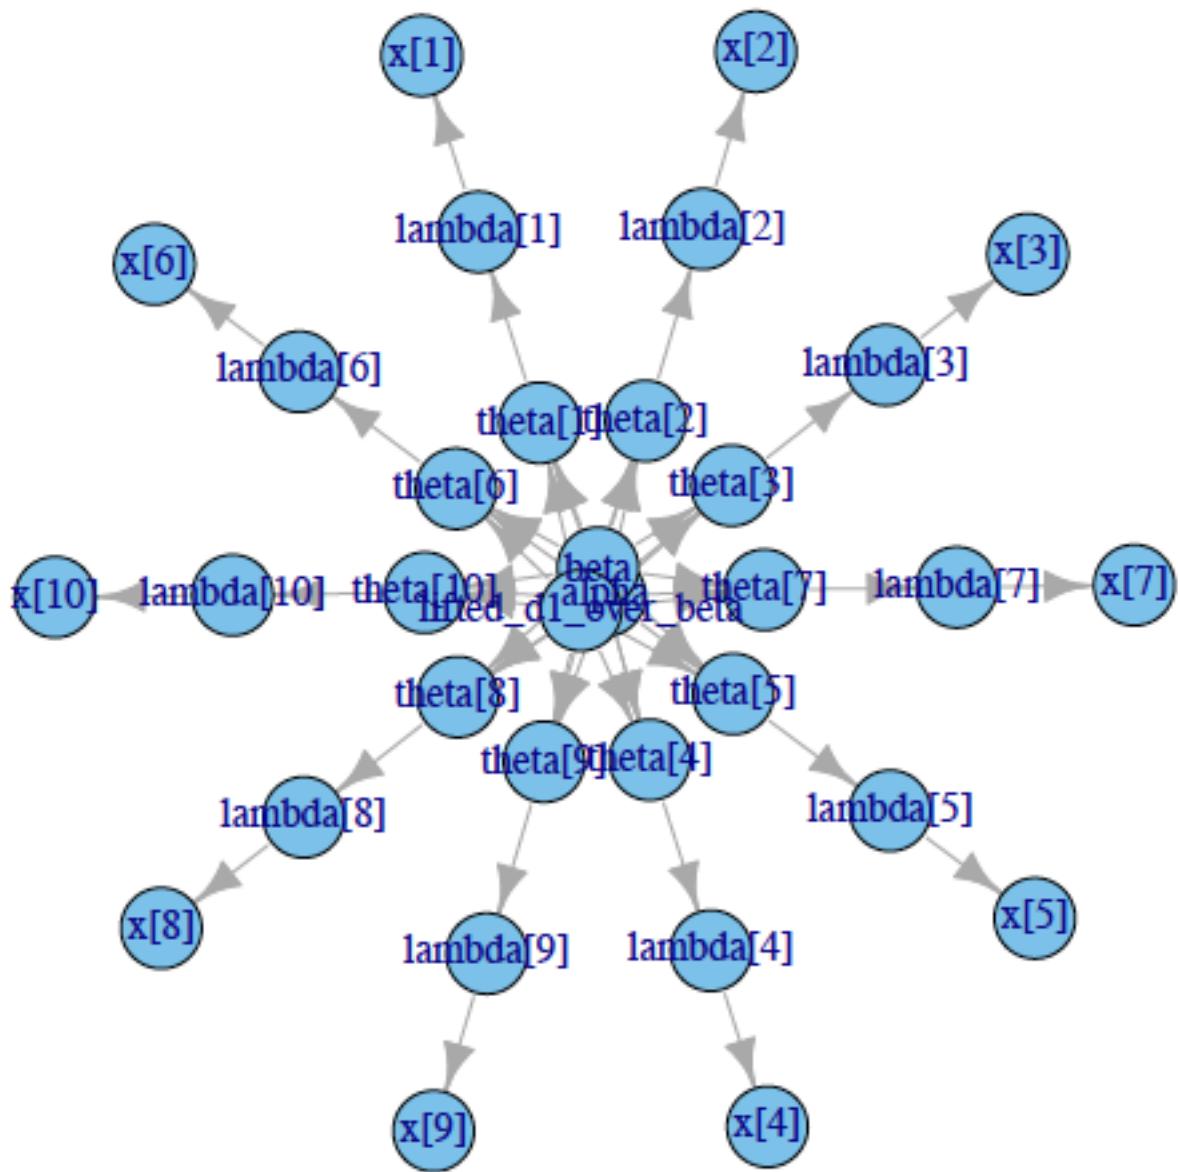


Figure due to igraph plot function

What NIMBLE does with BUGS models

We turn:

```
model {  
  for(i in 1:N) {  
    theta[i] ~ dgamma(alpha, beta)  
    lambda[i] <- theta[i] * tt[i]  
    x[i] ~ dpois(lambda[i])  
  }  
  alpha ~ dexp(1.0)  
  beta ~ dgamma(2, 2)  
}
```

into

```
> pumpModel$alpha <- 5  
  
> simulate(pumpModel, 'beta')  
  
> calculate(pumpModel,  
           c('lambda[1:5]', 'x[1:5]'))  
  
> calculate(pumpModel,  
           pumpModel$getDependencies( 'beta'))  
  
> getLogProb(pumpModel, 'x')
```

What NIMBLE does with BUGS models

We turn:

```
model {  
  for(i in 1:N) {  
    theta[i] ~ dgamma(alpha, beta)  
    lambda[i] <- theta[i] * tt[i]  
    x[i] ~ dpois(lambda[i])  
  }  
  alpha ~ dexp(1.0)  
  beta ~ dgamma(2, 2)  
}
```

into

```
> pumpModel$alpha <- 5  
  
> simulate(pumpModel, 'beta')  
  
> calculate(pumpModel,  
           c('lambda[1:5]', 'x[1:5]'))  
  
> calculate(pumpModel,  
           pumpModel$getDependencies( 'beta'))  
  
> getLogProb(pumpModel, 'x')
```

Use model variables

What NIMBLE does with BUGS models

We turn:

```
model {  
  for(i in 1:N) {  
    theta[i] ~ dgamma(alpha, beta)  
    lambda[i] <- theta[i] * tt[i]  
    x[i] ~ dpois(lambda[i])  
  }  
  alpha ~ dexp(1.0)  
  beta ~ dgamma(2, 2)  
}
```

into

```
> pumpModel$alpha <- 5  
  
> simulate(pumpModel, 'beta')  
  
> calculate(pumpModel,  
           c('lambda[1:5]', 'x[1:5]'))  
  
> calculate(pumpModel,  
           pumpModel$getDependencies( 'beta'))  
  
> getLogProb(pumpModel, 'x')
```

Simulate part of the model

What NIMBLE does with BUGS models

We turn:

```
model {  
  for(i in 1:N) {  
    theta[i] ~ dgamma(alpha, beta)  
    lambda[i] <- theta[i] * tt[i]  
    x[i] ~ dpois(lambda[i])  
  }  
  alpha ~ dexp(1.0)  
  beta ~ dgamma(2, 2)  
}
```

into

```
> pumpModel$alpha <- 5  
  
> simulate(pumpModel, 'beta')  
  
> calculate(pumpModel,  
           c('lambda[1:5]', 'x[1:5]'))  
  
> calculate(pumpModel,  
           pumpModel$getDependencies( 'beta'))  
  
> getLogProb(pumpModel, 'x')
```

Calculate log probabilities from part of the model

What NIMBLE does with BUGS models

We turn:

```
model {  
  for(i in 1:N) {  
    theta[i] ~ dgamma(alpha, beta)  
    lambda[i] <- theta[i] * tt[i]  
    x[i] ~ dpois(lambda[i])  
  }  
  alpha ~ dexp(1.0)  
  beta ~ dgamma(2, 2)  
}
```

into

```
> pumpModel$alpha <- 5  
  
> simulate(pumpModel, 'beta')  
  
> calculate(pumpModel,  
           c('lambda[1:5]', 'x[1:5]'))  
  
> calculate(pumpModel,  
           pumpModel$getDependencies( 'beta'))  
  
> getLogProb(pumpModel, 'x')
```

Query the model's structure:
How are variables connected?

What NIMBLE does with BUGS models

We turn:

```
model {  
  for(i in 1:N) {  
    theta[i] ~ dgamma(alpha, beta)  
    lambda[i] <- theta[i] * tt[i]  
    x[i] ~ dpois(lambda[i])  
  }  
  alpha ~ dexp(1.0)  
  beta ~ dgamma(2, 2)  
}
```

into

```
> pumpModel$alpha <- 5  
  
> simulate(pumpModel, 'beta')  
  
> calculate(pumpModel,  
           c('lambda[1:5]', 'x[1:5]'))  
  
> calculate(pumpModel,  
           pumpModel$getDependencies( 'beta'))  
  
> getLogProb(pumpModel, 'x')
```

Query previously calculated log probabilities

NIMBLE

1. BUGS language → R model object
2. NIMBLE programming language
 - ✧ Domain-specific language (DSL) embedded in R
 - ✧ Compilation via C++
3. Algorithm library: MLE, MCMC etc.

NIMBLE: Programming With Models

We want:

- High-level processing (model structure) in R
- Low-level processing in C++

NIMBLE: Programming With Models

- Say we want an objective function
 - Input: Values for (alpha, beta)
 - Output: Log probability density of part of the model that depends on (alpha, beta)
 - The “Markov blanket”
 - The we can optimize for:
 - Maximum likelihood estimation
 - 1st step of Laplace approximation / Quadrature
 - MCMC proposal density

NIMBLE: Programming With Models

```
objectiveFunction <- nimbleFunction (   
  setup = function(model, nodes) {  
    calcNodes <- model$getDependencies(nodes)  
  },  
  run = function(P = double(1)) {  
    values(model, nodes) <<- P  
    sumLogProb <- calculate(model, calcNodes)  
    return(sumLogProb)  
    returnType(double())  
  })
```

2 kinds of functions

NIMBLE: Programming With Models

```
objectiveFunction <- nimbleFunction ( pumpModel, c('alpha', 'beta')  
  setup = function(model, nodes) {  
    calcNodes <- model$getDependencies(nodes)  
  },  
  run = function(P = double(1)) {  
    values(model, nodes) <<- P  
    sumLogProb <- calculate(model, calcNodes)  
    return(sumLogProb)  
    returnType(double())  
  })
```

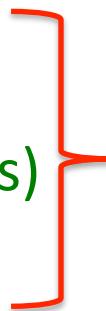
The diagram illustrates the two types of functions used in the code. A green bracket spans from the opening parenthesis of the first argument of the nimbleFunction call to the closing parenthesis of the run function definition. A red bracket groups the 'nimbleFunction' call and the 'function' definitions for 'setup' and 'run'. The text '2 kinds of functions' is written in red next to the red bracket.

NIMBLE: Programming With Models

```
objectiveFunction <- nimbleFunction (
```

```
  setup = function(model, nodes) {  
    calcNodes <- model$getDependencies(nodes)  
  },
```

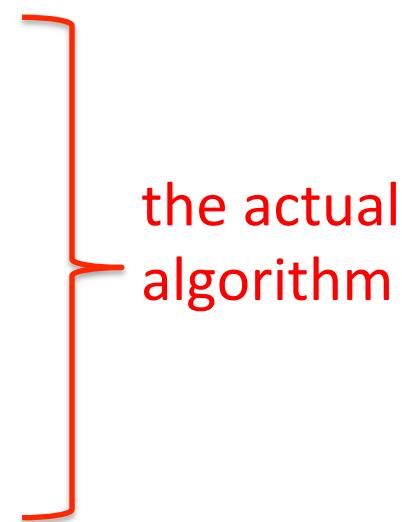
```
  run = function(P = double(1)) {  
    values(model, nodes) <<- P  
    sumLogProb <- calculate(model, calcNodes)  
    return(sumLogProb)  
    returnType(double())  
  })
```



query model
structure ONCE.

NIMBLE: Programming With Models

```
objectiveFunction <- nimbleFunction (   
  
  setup = function(model, nodes) {  
    calcNodes <- model$getDependencies(nodes)  
  },  
  
  run = function(P = double(1)) {  
    values(model, nodes) <<- P  
    sumLogProb <- calculate(model, calcNodes)  
    return(sumLogProb)  
    returnType(double())  
  })
```



the actual algorithm

Using a NIMBLE function

```
> objFun1 <- objectiveFunction(pumpModel,
                                c('alpha', 'beta'))
> objFun1$run(c(.9, 1.2))
-16.1

> CobjFun1 <- compileNimble(objFun1,
                               project = pumpModel)
> CobjFun1$run(c(0.9, 1.2))
-16.1
```

Using a NIMBLE function

```
> objFun1 <- objectiveFunction(pumpModel,
                                c('alpha', 'beta'))
> objFun1$run(c(.9, 1.2))
-16.1

> CobjFun1 <- compileNimble(objFun1,
                               project = pumpModel)
> CobjFun1$run(c(0.9, 1.2))
-16.1
```

Using a NIMBLE function

```
> objFun1 <- objectiveFunction(pumpModel,
                                c('alpha', 'beta'))
> objFun1$run(c(.9, 1.2))
-16.1

> CobjFun1 <- compileNimble(objFun1,
                               project = pumpModel)
> CobjFun1$run(c(0.9, 1.2))
-16.1
```

The NIMBLE compiler

One-line example:

nimbleFunction run code:

```
Y <- t(X) %*% X
```

The NIMBLE compiler

One-line example:

nimbleFunction run code:

`Y <- t(X) %*% X`



The NIMBLE compiler

One-line example:

nimbleFunction run code:

`Y <- t(X) %*% X`

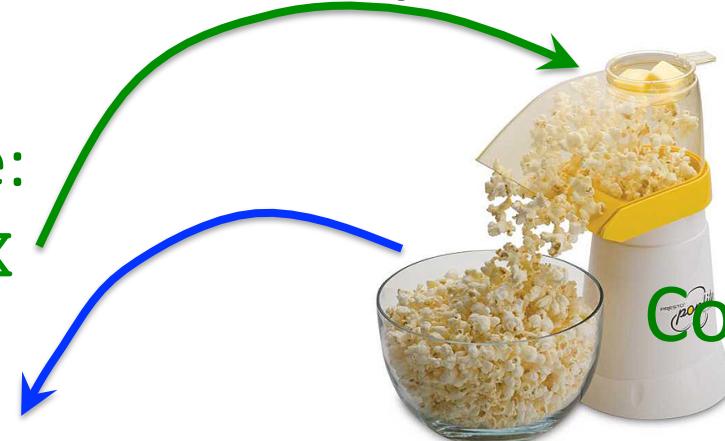
C++:

type declaration for Y, EigY, EigX

set size of Y

Point EigY at Y and EigX at X

`EigY = EigX.transpose() * EigX;`



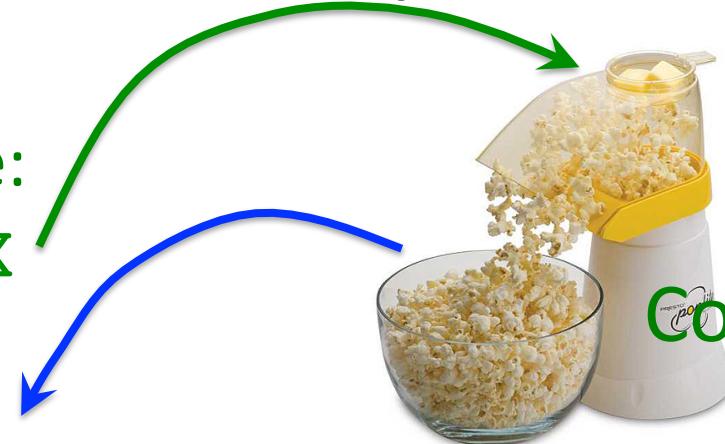
Uses Eigen
C++
library

The NIMBLE compiler

One-line example:

nimbleFunction run code:

```
Y <- t(X) %*% X
```



C++:

```
type declaration for Y, EigY, EigX  
set size of Y  
Point EigY at Y and EigX at X  
EigY = EigX.transpose() * EigX;
```

Uses Eigen
C++
library

R: compile the C++

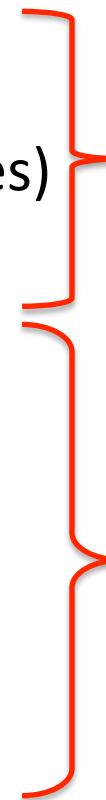
load the shared object

generate an R interface object (function or class)

NIMBLE compiler writes C++ classes

```
objectiveFunction <- nimbleFunction (
```

```
  setup = function(model, nodes) {  
    calcNodes <- model$getDependencies(nodes)  
  },  
  
  run = function(P = double(1)) {  
    values(model, nodes) <<- P  
    sumLogProb <- calculate(model, calcNodes)  
    return(sumLogProb)  
    returnType(double())  
  })
```



Member data

Member function(s)

The NIMBLE compiler

Feature summary:

- R-like matrix algebra (using Eigen library)
- R-like indexing (e.g. $X[1:5,]$)
- Use of model variables and nodes
- Sequential integer iteration
- if-then-else, do-while
- Declare input & output types only
- Access to much of Rmath.h (e.g. distributions)
- Automatic R interface / wrapper
- Many improvements / extensions planned
- Please contribute!

Keeping track of sets of model variables: modelValues

```
> mv <- modelValues(pumpModel, 1000)
```

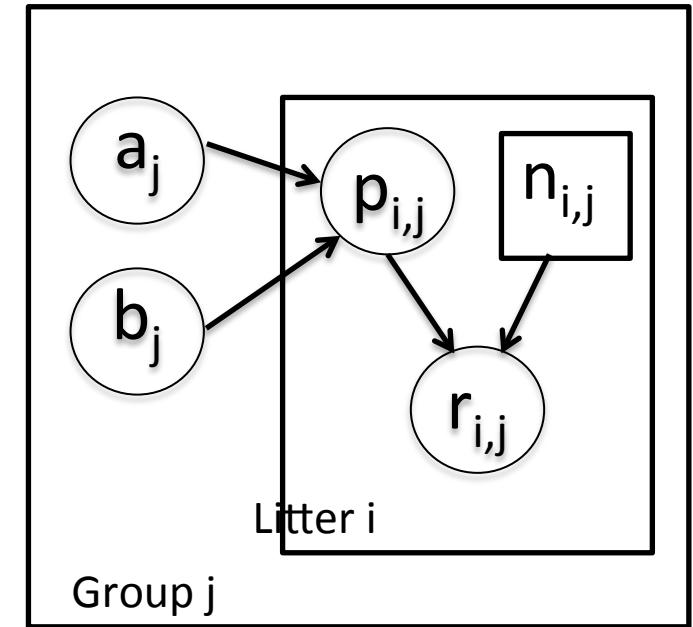
```
> mv['alpha', 3] <- 5
```

```
> copy(model, mv, rowTo = 10)
```

NIMBLE MCMC: the Litters Example

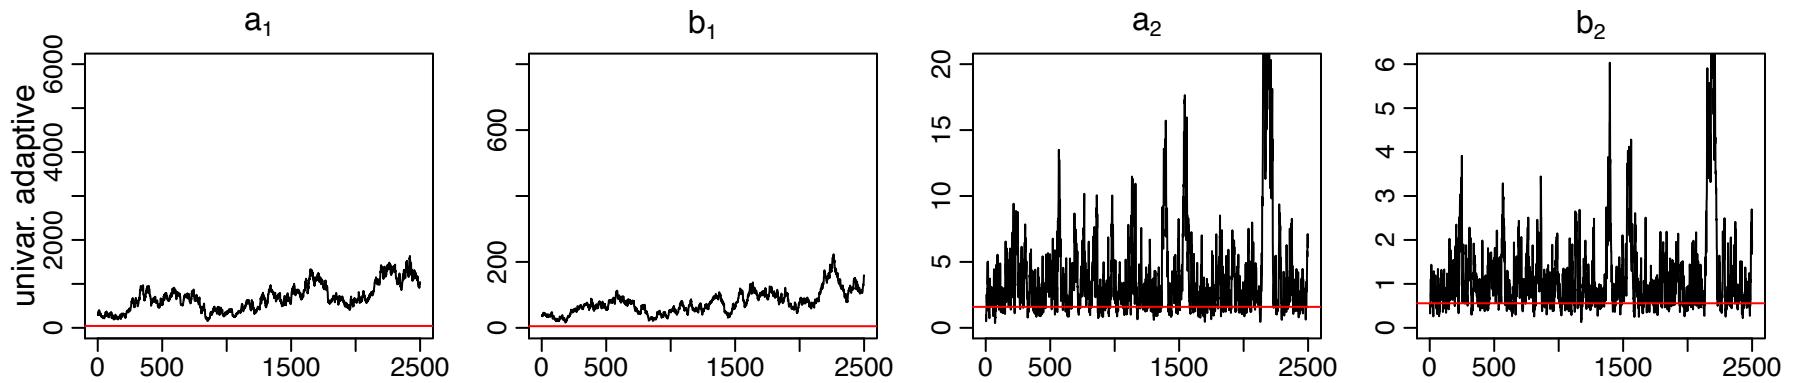
Beta-binomial for clustered binary response data

```
littersCode <- modelCode({  
  for(j in 1:G) {  
    for(l in 1:N) {  
      r[i, j] ~ dbin(p[i, j], n[i, j]);  
      p[i, j] ~ dbeta(a[j], b[j]);  
    }  
    mu[j] <- a[j]/(a[j] + b[j]);  
    theta[j] <- 1.0/(a[j] + b[j]);  
    a[j] ~ dgamma(1, 0.001);  
    b[j] ~ dgamma(1, 0.001);  
  }  
})
```



Default MCMC: Gibbs + Metropolis

```
> littersMCMCspec <- MCMCspec(littersModel,  
                                list(adaptInterval = 100))  
  
> littersMCMC <- buildMCMC(littersMCMCspec)  
  
> littersMCMC_cpp <- compileNIMBLE(littersModel,  
                                       project = littersModel)  
  
> littersMCMC_cpp(10000)
```



Blocked MCMC: Gibbs + Blocked Metropolis

```
> littersMCMCspec2 <- MCMCspec(littersModel, list(adaptInterval = 100))

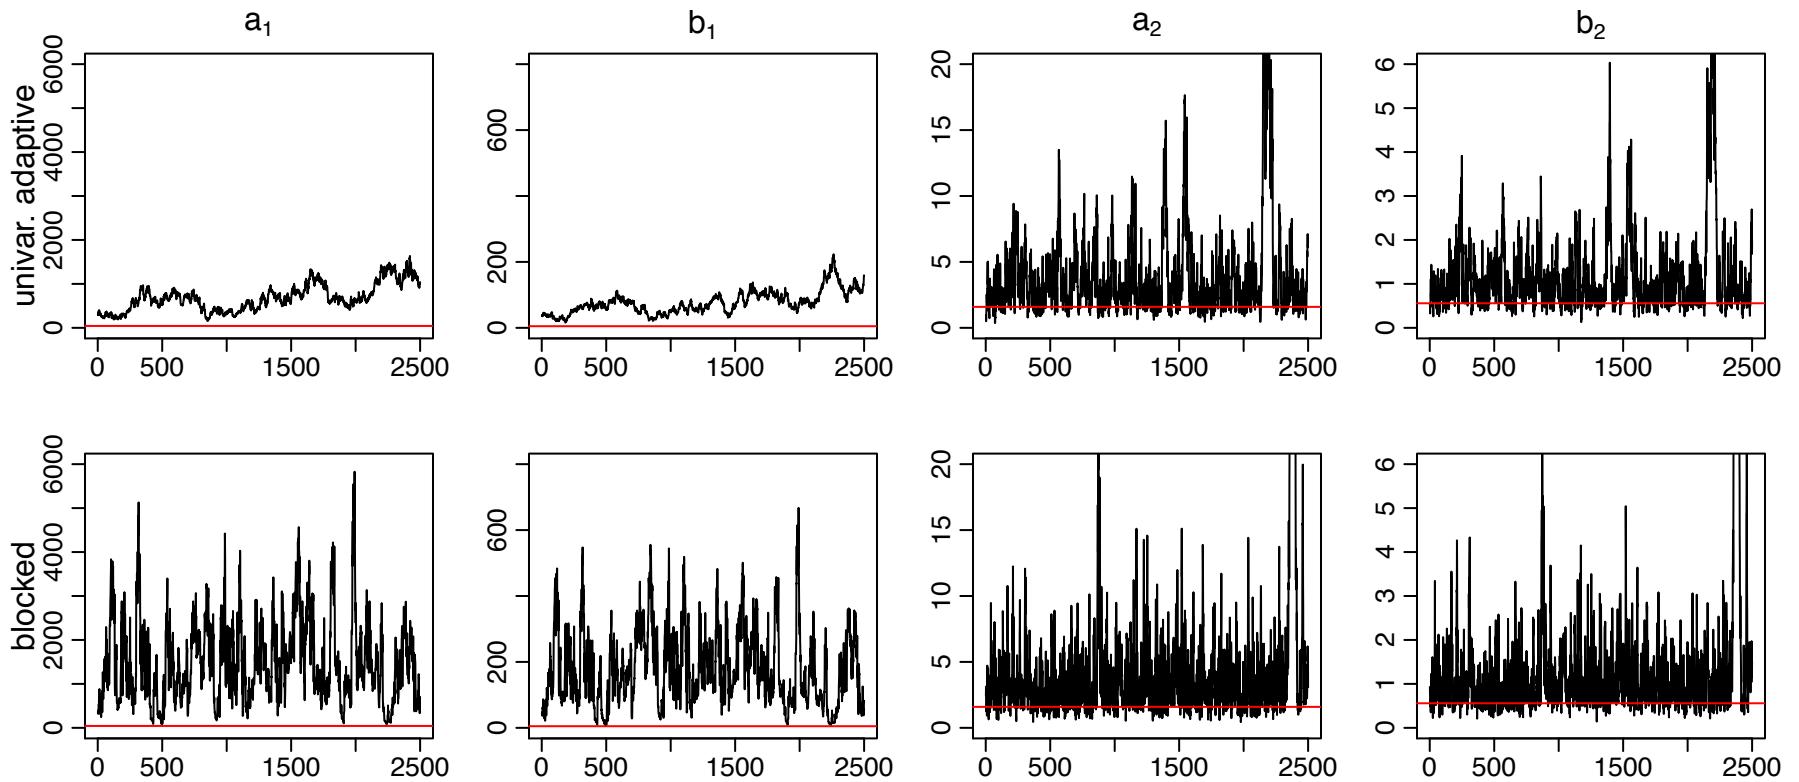
> littersMCMCspec2$addSampler('RW_block', list(targetNodes = c('a[1]', 'b[1]'),
adaptInterval = 100)

> littersMCMCspec2$addSampler('RW_block', list(targetNodes = c('a[2]', 'b[2]'),
adaptInterval = 100)

> littersMCMC2 <- buildMCMC(littersMCMCspec2)

> littersMCMC2_cpp <- compileNIMBLE(littersMCMC2, project = littersModel)

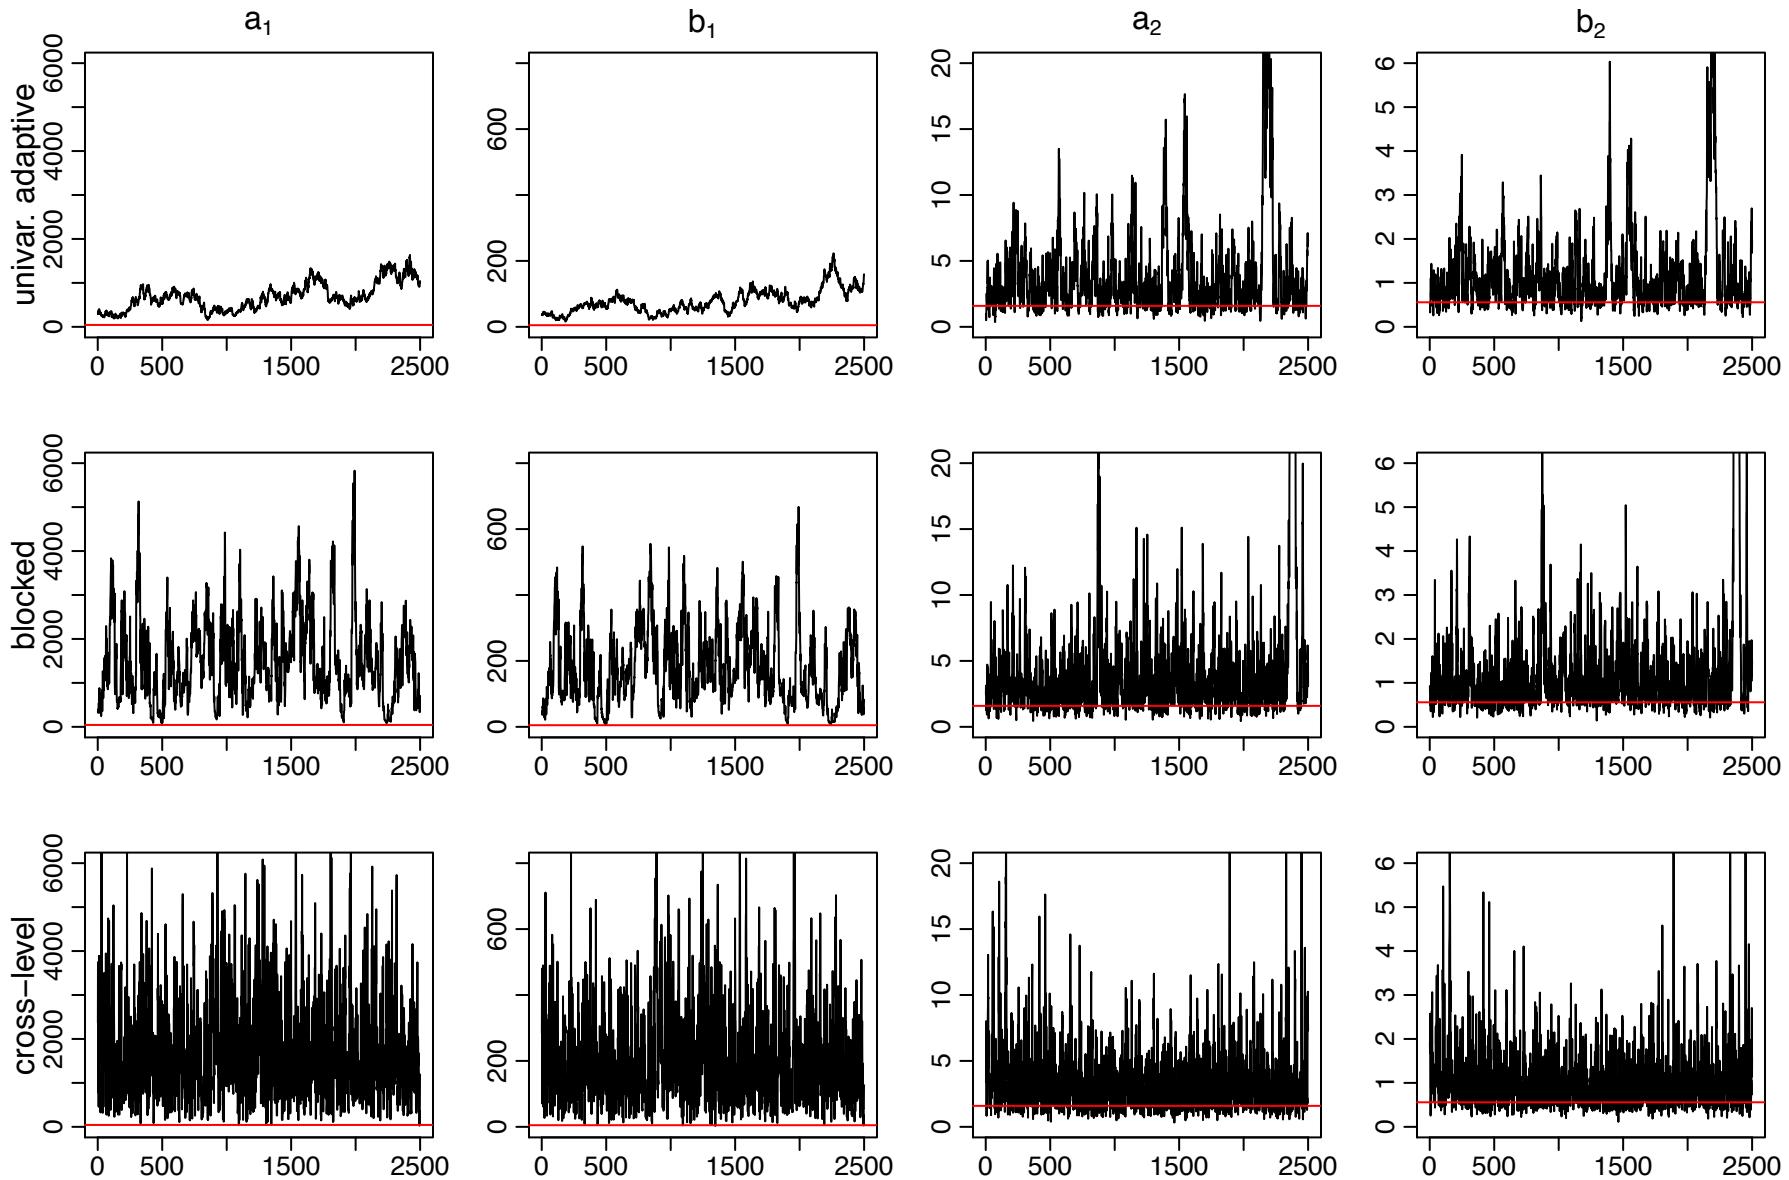
> littersMCMC2_cpp(10000)
```



Cross-level MCMC: Gibbs + Cross-level Updaters

- Problem: Correlation across levels of model
- New “cross-level” updater in NIMBLE language.
 - Random walk on one level
 - Joint conjugate proposal on dependent level

```
...
> littersMCMCspec3$addSampler(
  'crossLevel',
  list(topNodes = c('a[1]', 'b[1]'),
       adaptInterval = 100) )
...
...
```



<http://R-nimble.org>

